

pyspark Practice Notes

content:-

Read and Write files using
PySpark

PySpark show()

Run SQL Queries with PySpark

PySpark Pandas API

Select columns in PySpark
dataframe

PySpark withColumn()

PySpark Drop Columns

PySpark Rename Columns

PySpark Filter vs Where

PySpark orderBy() and sort()

PySpark GroupBy()

PySpark Pivot

PySpark Joins

PySpark Union

Spark Session -

SparkSession is designed to be a singleton, which means that only one instance should be active in the application at any given time.

```
spark2 = SparkSession.newSession  
  
print(spark2)  
  
spark = SparkSession.builder \  
    .appName("My PySpark Application") \  
    .master("local[*]") \  
    .config("spark.executor.memory", "4g") \  
    .config("spark.executor.cores", "4") \  
    .config("spark.driver.memory", "2g") \  
    .getOrCreate()
```

What is SparkSession – PySpark Entry Point, Dive into SparkSession

Jagdeesh

Introduction

PySpark, the Python library for Apache Spark, has gained immense popularity among data engineers and data scientists due to its simplicity and power in handling big data tasks.

This blog post will provide a comprehensive understanding of the PySpark entry point, the SparkSession. We'll explore the concepts, features, and the use of SparkSession to set up a PySpark application effectively.

What is SparkSession?

SparkSession is the entry point for any PySpark application, introduced in Spark 2.0 as a unified API to replace the need for separate SparkContext, SQLContext, and HiveContext.

The `SparkSession` is responsible for coordinating various Spark functionalities and provides a simple way to interact with structured and semi-structured data, such as reading and writing data from various formats, executing SQL queries, and utilizing built-in functions for data manipulation.

`SparkSession` offers several benefits that make it an essential component of PySpark applications:

Simplified API: `SparkSession` unifies the APIs of `SparkContext`, `SQLContext`, and `HiveContext`, making it easier for developers to interact with Spark's core features without switching between multiple contexts.

Configuration management: You can easily configure a `SparkSession` by setting various options, such as the application name, the master URL, and other configurations.

Access to Spark ecosystem: `SparkSession` allows you to interact with the broader Spark ecosystem, such as `DataFrames`, `Datasets`, and `MLlib`, enabling you to build powerful data processing pipelines.

Improved code readability: By encapsulating multiple Spark contexts, `SparkSession` helps you write cleaner and more maintainable code.

Creating a `SparkSession`

To create a `SparkSession`, we first need to import the necessary PySpark modules and classes. Here's a simple example:

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \  
    .appName("My PySpark Application") \  
    .master("local[*]") \  
    .getOrCreate()
```

In this example, we import the `SparkSession` class from the `pyspark.sql` module and use the `builder` method to configure the application name and master URL. The `getOrCreate()` method is then used to either get the existing `SparkSession` or create a new one if none exists.

The `SparkSession.builder` object provides various functions to configure the `SparkSession` before creating it. Some of the important functions are:

`appName(name)`: Sets the application name, which will be displayed in the Spark web user interface.

`master(masterURL)`: Sets the URL of the cluster manager (like YARN, Mesos, or standalone) that Spark will connect to. You can also set it to "local" or "local[N]" (where N is the number of cores) for running Spark locally.

`config(key, value)`: Sets a configuration property with the specified key and value. You can use this method multiple times to set multiple configuration properties.

`config(conf)`: Sets the Spark configuration object (`SparkConf`) to be used for building the `SparkSession`.

`enableHiveSupport()`: Enables Hive support, including connectivity to a persistent Hive metastore, support for Hive SerDes, and Hive user-defined functions (UDFs).

`getOrCreate()`: Retrieves an existing `SparkSession` or, if there is none, creates a new one based on the options set via the builder.

How many pyspark sessions can be created?

In PySpark, you can technically create multiple `SparkSession` instances, but it is not recommended. The standard practice is to use a single `SparkSession` per application.

`SparkSession` is designed to be a singleton, which means that only one instance should be active in the application at any given time

Create new `SparkSession`

```
spark2 = SparkSession.newSession
```

```
print(spark2)
```

Configuring a `SparkSession`

You can configure a `SparkSession` with various settings, such as the number of executor cores, executor memory, driver memory, and more. Here's an example

```
spark = SparkSession.builder \
    .appName("My PySpark Application") \
    .master("local[*]") \
    .config("spark.executor.memory", "4g") \
    .config("spark.executor.cores", "4") \
    .config("spark.driver.memory", "2g") \
    .getOrCreate()
```

In this example, we've added three additional configurations for executor memory, executor cores, and driver memory using the `config()` method.

Accessing SparkSession Components

```
# Access SparkContext
```

```
spark_context = spark.sparkContext
```

```
# Access SQLContext
```

```
sql_context = spark._wrapped
```

```
# Read a CSV file
```

```
data_frame = spark.read.csv("path/to/your/csv-file", header=True, inferSchema=True)
```

```
# Write the data to a Parquet file
```

```
data_frame.write.parquet("path/to/output/parquet-file")
```

Executing SQL Queries with SparkSession

With SparkSession, you can also execute SQL queries directly on your data. Here's an example:

```
# Register a DataFrame as a temporary table
```

```

data_frame.createOrReplaceTempView("my_table")

# Execute an SQL query on the temporary table
result = spark.sql("SELECT * FROM my_table WHERE age > 30")

# Display the result
result.show()

```

Example: Word Count with SparkSession

```

from pyspark.sql import SparkSession

from pyspark.sql.functions import split, explode, col

# creating a spark session
spark = SparkSession.builder \
    .appName("Counting ") \
    .master("local[*]") \
    .getOrCreate()

# read the input file
data = spark.read.txt("txt file")

# Split the line into words
words = data.select(explode(split(col("Value", " "))).alias("Word"))

```

```
# perform the word count

word_count = words.groupBy("Word").count()

# Display the results

word_count.show()

spark.stop()
```

Note -

explode - EXPLODE is a PySpark function used to work over columns in PySpark.

EXPLODE is used for the analysis of nested column data.

PySpark EXPLODE converts the Array of Array Columns to row.

EXPLODE can be flattened up post analysis using the flatten method.

EXPLODE returns type is generally a new row for each element given.

explode() will return each and every individual value from an array. If the array is empty or null, it will ignore and go to the next array in an array type column in PySpark DataFrame. This is possible using the select() method. Inside this method, we can use the array_min() function and return the result.

Read and Write files using PySpark – Multiple ways to Read and Write data using PySpark

```
pip install findspark

pip install pyspark

import findspark

findspark.init()
```



```

from pyspark.sql import SparkSession

spark = SparkSession.builder \

    .appName("My name is ") \

    .master("Local[*]") \

    .getOrCreate()

```

Creating a Pyspark dataframe

```

data = [("Alice", 34), ("Bob", 45), ("Cathy", 29)]

columns = ["Name", "Roll_Number"]

df = spark.createDataFrame(data, columns)

df.show()

```

```

+-----+----+
| Name|Age|
+-----+----+
|Alice| 34|
| Bob| 45|
|Cathy| 29|

```

--	--

```

csv_file = "path/to/your/csv/file.csv"

```

```
df_csv = spark.read.csv(csv_file, header=True, inferSchema=True)
```

```
output_path = "path/to/output/csv/file.csv"
```

```
df_csv.write.csv(output_path, header=True, mode="overwrite")
```

```
json_file = "path/to/your/json/file.json"
```

```
df_json = spark.read.json(json_file)
```

```
output_path = "path/to/output/json/file.json"
```

```
df_json.write.json(output_path, mode="overwrite")
```

Reading and Writing Parquet Files

To read a Parquet file using PySpark, you can use the `read.parquet()` method:

```
parquet_file = "path/to/your/parquet/file.parquet"
```

```
df_parquet = spark.read.parquet(parquet_file)
```

To write the data back to a Parquet file, use the `write.parquet()` method:

```
output_path = "path/to/output/parquet/file.parquet"
```

```
df_parquet.write.parquet(output_path, mode="overwrite")
```

```
data [  
    {"Name": "Alice", "Age": 30, "City": "New York"},  
    {"Name": "Bob", "Age": 25, "City": "San Francisco"},  
    {"Name": "Charlie", "Age": 35, "City": "Los Angeles"}]  
  
df = spark.createDataFrame(data)  
  
df.createOrReplaceTempView("People")  
  
query = "select * from People where Age > 30"  
  
result_df = spark.sql(query)  
  
result_df.show()
```

Convert Pandas Dataframe to PySpark Dataframe

```
import pandas as pd  
  
data [  
    {"Name": "Alice", "Age": 30, "City": "New York"},  
    {"Name": "Bob", "Age": 25, "City": "San Francisco"},  
    {"Name": "Charlie", "Age": 35, "City": "Los Angeles"}]  
  
pandasData = pd.DataFrame(data, columns = ["Name", "Age", "City"])
```

```
# creating the pyspark dataframe from pandas

sparkdf = spark.createDataFrame(pandasData)

sparkdf.printSchema()

sparkdf.show()
```

root

```
-- name: string (nullable = true)
-- age: long (nullable = true)
-- city: string (nullable = true)
```

```
+-----+---+-----+
|  name|age|    city|
+-----+---+-----+
| Alice| 30| New York|
|  Bob| 25|San Francisco|
|Charlie| 35| Los Angeles|
```

--	--	--

A DataFrame is a distributed collection of data organized into named columns.

It is conceptually equivalent to a table in a relational database or a data frame in Python, but optimized for large-scale processing. DataFrame as a spreadsheet with rows and columns.

The `show()` function is a method available for DataFrames in PySpark. It is used to display the contents of a DataFrame in a tabular format, making it easier to visualize and understand the data.

Syntax

`DataFrame.show(n=20, truncate=True, vertical=False)`

Parameters:

`n`: The number of rows to display. The default value is 20.

`truncate`: If set to True, the column content will be truncated if it is too long. The default value is True.

`vertical`: If set to True, the output will be displayed vertically. The default value is False.

Example 1-

```
import findaspark

findspark.init()

from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName \
    .getOrCreate()

data = [("Alice":34 ), ("Bob":45), ("Charlie":29), ("David": 31)]

columns = ["Name", "Age"]

df = spark.createDataFrame(data, columns)

df.show()
```

```
+-----+---+
|  Name|Age|
+-----+---+
| Alice| 34|
|  Bob| 45|
|Charlie| 29|
|  David| 31|
```

Display the Specific number of row. like 2

```
df.show(2)

+-----+---+
|  Name|Age|
+-----+---+
|Alice| 34|
|  Bob| 45|
+-----+---+

only showing top 2 rows
```

Display Contents Without Truncation

```
df.show(truncate = False)

Name   |Age|
+-----+---+
|Alice |34 |
```

|Bob |45 |

|Charlie|29 |

|David |31 |

Display Contents Vertically

`df.show(vertical=True)`

-RECORD 0-----

Name | Alice

Age | 34

-RECORD 1-----

Name | Bob

Age | 45

-RECORD 2-----

Name | Charlie

Age | 29

-RECORD 3-----

Name | David

Age | 31

Run SQL Queries with PySpark

OrderID	ProductID	Quantity	Price	OrderDate
---------	-----------	----------	-------	-----------

1	101	3	100	2023-01-01
2	102	1	200	2023-01-02
3	101	2	100	2023-01-03
4	103	5	50	2023-01-04

name of file = sales data

```
csv_file = sales_data.csv
```

```
df = spark.read.option("header":"true").option("inferSchema":"true").csv(csv_files)
```

```
# creatting a temporary View
```

```
df.createOrReplaceTempView("Saless_data")
```

Calculating the total Revenue of each product

```
query = """
```

```
select ProductID,
```

```
SUM (Quantity * Price) AS totalRevenue
```

```
from saless_data
```

```
group by ProductID
```

```
"""
```

```
result = spark.sql(query)
```

```
result.show()
```



```

+-----+-----+
|ProductID|TotalRevenue|
+-----+-----+
| 101 | 500 |
| 102 | 200 |
| 103 | 250 |
+-----+-----+

```

to find the top 2 products with the highest revenue

```

query = """
select
ProductID,
SUM (Quantity * Price) AS totalRevenue
from saless_data
group by ProductID
order by totalRevenue
limit 2; """
resultDF = spark.sql(query)
resultDF.show()

```

```

+-----+-----+

```

ProductID	TotalRevenue
101	500
102	200

PySpark Pandas API

```
import pandas as pd
```

```
import numpy as np
```

```
from pyspark.sql import SparkSession
```

```
import databricks.koalas as ks
```

```
# creatting the sparkSession
```

```
spark = SparkSession.builder \
```

```
    .appName \
```

```
    .getOrCreate()
```

reading the file

```
sales_data = ks.read_csv("sales_data.csv")
```

calculate the average revenue per unit sold and add it as a new column

```
sales_data['Avg_Revenue_per_Unit'] =  
sales_data['Revenue'] / sales_data['Unit_Sold']
```

Computing summary statistics

```
summary_stats = sales_data.groupby(['Store_ID', 'Product_ID']).agg(  
    {'Revenue': 'sum', 'Units_Sold': 'sum'}).reset_index()
```

Sorting the results

```
sorted_summary_stats = summary_stats.sort_values(  
    by=['Store_ID', 'Revenue'], ascending=[True, False])
```

Exporting the results

```
sorted_summary_stats.to_csv("summary_stats.csv", index=False)
```

```
spark.stop()    # for stoping that current session.
```

Select columns in PySpark dataframe

```
import findspark  
  
findspark.init()  
  
spark = SparkSession.builder \  
    .appName \  
    .master \  
    .getOrCreate()
```

```
data = [("Alice", 34, "Female"), ("Bob":45, "Male"), ("Charlie":28, "Male"),
("Diana", 39, "Female")]

columns = ["Name", "Age", "Gender"]

df = spark.createDataFrame(data, columns)

df.show()
```

```
  Name|Age|Gender|
+-----+---+-----+
| Alice| 34|Female|
|  Bob| 45|  Male|
|Charlie| 28|  Male|
| Diana| 39|Female|
+-----+---+-----+
```

Select columns using column names

```
selectdf1 = df.select("Name", "Age")

selectdf1.show()
```

```
  Name|Age|
+-----+---+
| Alice| 34|
|  Bob| 45|
|Charlie| 28|
| Diana| 39|
```

select column using the col functions

```
selectdf2 = df.select (col("name"), col("Age"))
```

```
selectdf2.show()
```

```
Name|Age|
```

```
+-----+---+
```

```
| Alice| 34|
```

```
|  Bob| 45|
```

```
|Charlie| 28|
```

```
| Diana| 39
```

Selecting Columns using the '[']' Operator

```
# Select a single column using the '[' operator
```

```
name_df = df["Name"]
```

```
# Select multiple columns using the '[' operator
```

```
selected_df3 = df.select(df["Name"], df["Age"])
```

```
selected_df3.show()
```

```
Name|Age|
```

```
+-----+---+
```

```
| Alice| 34|
```

```
|  Bob| 45|
```

```
|Charlie| 28|
```

```
| Diana| 39|
```

Select Columns using index

```
# Define the column indices you want to select
```

```
column_indices = [0, 2]
```

```
# Extract column names based on indices
```

```
selected_columns = [df.columns[i] for i in column_indices]
```

```
# Select columns using extracted column names
```

```
selected_df4 = df.select(selected_columns)
```

```
# Show the result DataFrame
```

```
selected_df4.show()
```

```
Name|Gender|
```

```
+-----+-----+
```

```
| Alice|Female|
```

```
|  Bob|  Male|
```

```
|Charlie|  Male|
```

```
| Diana|Female|
```

--	--

Selecting Columns using the 'withColumn' and 'drop' Functions.

select specific columns while adding or removing columns, you can use the **‘withColumn’** function to add a new column and the **‘drop’** function to remove a column.

```
# Add a new column 'IsAdult' and remove the 'Gender' column
```

```
selectDF = df.withColumn("IsAdult", col("Age") > 18).drop("Gender")  
selectDF.show()
```

Name	Age	IsAdult
Alice	34	true
Bob	45	true
Charlie	28	true
Diana	39	true

Selecting Columns using SQL Expressions

select columns using the **‘selectExpr’** function. This is useful when you want to perform operations on columns while selecting them.

```
selectDF = df.selectExpr("Name", "Age", "Age >= 18 AS IsAdult")  
selectDF.show()
```

Name	Age	IsAdult
------	-----	---------

```
+-----+---+-----+
| Alice| 34|  true|
|  Bob| 45|  true|
|Charlie| 28|  true|
| Diana| 39|  true|
```

PySpark withColumn

The “**withColumn**” function in PySpark allows to **add, replace, or update** columns in a **DataFrame**. It is a **DataFrame** transformation operation, meaning it returns a new **DataFrame** with the specified changes, without altering the original **DataFrame**.

The “**withColumn**” function is particularly useful when we need to perform column-based operations like renaming, changing the data type, or applying a function to the values in a column.

Syntax

```
DataFrame.withColumn(colName, col)
```

where:

DataFrame: The original PySpark **DataFrame** you want to manipulate.

colName: The name of the new or existing column you want to add, replace, or update.

col: The new expression or value for the specified column.


```

import findspark

findspark.init()

from pyspark.sql import SparkSession

from pyspark.sql.functions import col

# initialization the sparksession

spark = SparkSession.builder \

    .appName \

    .getOrCreate()

data = [(1, "Alice", 25),

(2, "Bob", 30),

(3, "Charlie", 35)

]

columns = ["id", "name", "age"]

# create a Dataframe

df = spark.createDataFrame(data, columns)

df.show()

```

```

id|  name|age|
+---+-----+---+
| 1|  Alice| 25|
| 2|   Bob| 30|
| 3|Charlie| 35|

```

1. Remaining the column.

"withColumns" to rename the "age" column to "years"

rename the age column to years

```
df = df.withColumn("Years", col("age"))
```

drop the column age

```
df = df.drop("Age")
```

```
df.show()
```

```
id|  name|years|
```

```
+---+-----+-----+
```

```
| 1| Alice| 25|
```

```
| 2|  Bob| 30|
```

```
| 3|Charlie| 35|
```

```
+---+-----+-----
```

2. Applying a function to a column

Converting the age from years to months

```
from pyspark.sql.functions import expr

df = df.withColumn("months", expr("years * 12"))

df.show()
```

```
id|  name|years|months|
+---+-----+-----+-----+
| 1| Alice|  25|  300|
| 2|  Bob|  30|  360|
| 3|Charlie| 35|  420|
```

3. Change the column data type

```
from pyspark.sql.types import StringType

# change the id to string

df = df.withColumn("id", col("id").cast(StringType()))

df.show()
```

```
id|  name|years|months|
+---+-----+-----+-----+
| 1| Alice|  25|  300|
| 2|  Bob|  30|  360|
| 3|Charlie| 35|  420|
```

4. Conditional column update with "withColumn".

```

from pyspark.sql.functions import when

data = [(1, "Alice", 25, 45000), (2, "Bob", 30, 55000), (3, "Charlie", 35, 60000)]

columns = ["id", "name", "age", "salary"]

df = spark.createDataFrame(data, columns)

df.show()

# add the "tax" column based on the "salary".

df = df.withColumn("tax", when(col("salary")>=50000, col("salary")*0.1).otherwise(col("salary")
* 0.05

df.show()

```

id	name	age	salary
1	Alice	25	45000
2	Bob	30	55000
3	Charlie	35	60000

id	name	age	salary	tax
1	Alice	25	45000	2250.0
2	Bob	30	55000	5500.0
3	Charlie	35	60000	6000.0

5. Using the User Defined Function with withColumn.

```
from pyspark.sql.functions import udf
from pyspark.sql.types import StringType
```

```
def age_group(age):
    if age < 30:
        return "Young"
    elif age < 45:
        return "MiddleAged"
    else:
        return "Old"
```

```
# Register the UDF
```

```
age_group_udf = udf(age_group, StringType())
```

```
# add the new column age_group based on the age
```

```
df = df.withColumn("age_group", age_group_udf(col("age")))
```

```
df.show()
```

```
id|  name|age|salary|  tax| age_group|
```

```
+---+-----+---+-----+-----+-----+
```

```
| 1| Alice| 25| 45000|2250.0|   Young|
```

```
| 2|  Bob| 30| 55000|5500.0|Middle-aged|
```

```
| 3|Charlie| 35| 60000|6000.0|Middle-aged|
```

```
# creating the new column net_sales column based on the subtracting the tax column from salary
```

```
df = df.withColumn("net_sales", round(col("salary")- col("tax"), 2))
```

```
df.show()
```

```
id|  name|age|salary|  tax| age_group|net_salary|
```

```
+---+-----+---+-----+-----+-----+-----+
```

```
| 1| Alice| 25| 45000|2250.0|   Young| 42750.0|
```

```
| 2|  Bob| 30| 55000|5500.0|Middle-aged| 49500.0|
```

```
| 3|Charlie| 35| 60000|6000.0|Middle-aged| 54000.0|
```

```
# combining the multiple column into one column.
```

```
name and age_group column = name_age_group
```

We will use the “**concat_ws**” function, which allows us to concatenate multiple columns with a specified delimiter.

```
from pyspark.sql.functions import concat_ws
```

```
# combining the column
```

```
df = df.withColumn("name_age_group", concat_ws("_", col("name"), col("age_group")))
```

```
df.show()
```

id	name	age	salary	tax	age_group	net_salary	name_age_group
1	Alice	25	45000	2250.0	Young	42750.0	Alice - Young
2	Bob	30	55000	5500.0	Middle-aged	49500.0	Bob - Middle-aged
3	Charlie	35	60000	6000.0	Middle-aged	54000.0	Charlie - Middle-aged

PySpark Drop Columns

Drop() function to remove columns from a DataFrame

Creating the dataframe

```
import findspark

findspark.init()

from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName \
    .getOrCreate()

data = [("Alice", 30, "New York", "F"),
        ("Bob", 28, "San Francisco", "M"),
        ("Cathy", 29, "Los Angeles", "F"),
        ("David", 32, "Chicago", "M")]

```

```
columns= ["name", "age", "city", "gender"]

df = df.createDataFrame(data, columns)

df.show()
```

```
name|age|    city|gender|
+---+---+-----+-----+
|Alice| 30|  New York|    F|
| Bob| 28|San Francisco|    M|
|Cathy| 29| Los Angeles|    F|
|David| 32|   Chicago|    M|
```

Different ways to drop columns in PySpark DataFrame

Dropping a Single Column

The Drop() function can be used to remove a single column from a DataFrame.

Syntax:

```
df = df.drop("gender")

df.show()
```

```
name|    city|
+---+-----+
|Alice|  New York|
| Bob|San Francisco|
|Cathy| Los Angeles|
|David|   Chicago|
```

Dropping Multiple Columns

Drop() function to remove multiple columns from a DataFrame. Simply pass a list of column names to the function.

```
df = df.drop("age", "gender")  
  
df.show()
```

or list of column name

```
dropping_column_name = ["age", "gender"]  
  
df = df.drop(*dropping_column_name)  
  
df.show()
```

```
name|    city|  
+----+-----+  
|Alice| New York|  
| Bob|San Francisco|  
|Cathy| Los Angeles|  
|David|  Chicago
```

Dropping Columns Conditionally

```
if "gender" in df.columns:  
    df = df.drop("gender")  
  
df.show()  
  
name|age|    city|  
+----+---+-----+  
|Alice|25| New York|  
| Bob|30|San Francisco|  
|Cathy|28| Los Angeles|  
|David|35|  Chicago
```

```
|Alice| 30|   New York|
| Bob| 28|San Francisco|
|Cathy| 29| Los Angeles|
|David| 32|   Chicago|
```

Dropping Columns Using Regex Pattern

“drop()” function in combination with a regular expression (regex) pattern to drop multiple columns matching the pattern.

```
from pyspark.sql.functions import col
import re
regex_pattern = "gender|age"
df = df.select([col(c) for c in df.columns if not re.match(regex_pattern, c)])
df.show()
```

```
+---+-----+
| name|      city|
+---+-----+
|Alice|   New York|
| Bob|San Francisco|
|Cathy| Los Angeles|
|David|   Chicago|
```

PySpark Rename Columns

```
import findspark

findspark.init()

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("PySpark Rename Columns").getOrCreate()

from pyspark.sql import Row

data = [Row(name="Alice", age=25, city="New York"),
        Row(name="Bob", age=30, city="San Francisco"),
        Row(name="Cathy", age=35, city="Los Angeles")]

sample_df = spark.createDataFrame(data)

sample_df.show()
```

```
name|age|    city|
+---+---+-----+
|Alice| 25|  New York|
| Bob| 30|San Francisco|
|Cathy| 35| Los Angeles|
+---+---+-----+
```

Different ways to rename columns in a PySpark DataFrame

Renaming Columns Using 'withColumnRenamed'

The 'withColumnRenamed' method is a simple way to rename a single column in a DataFrame

```
renamed_df = sample_df.withColumnRenamed("age", "user_age")
```

```
renamed_df.show()
```

```
name|user_age|    city|
+----+-----+-----+
|Alice|    25|   New York|
| Bob|    30|San Francisco|
|Cathy|    35| Los Angeles
```

Renaming Columns Using 'select' and 'alias'

```
from pyspark.sql.functions import col

renamed_df = sample_df.select(col("name"), col("age").alias("user_age"), col("city"))

renamed_df.show()

name|user_age|    city|
+----+-----+-----+
|Alice|    25|   New York|
| Bob|    30|San Francisco|
|Cathy|    35| Los Angeles
```

Renaming Columns Using 'toDF'

```
renamed_df = sample_df.toDF("user_name", "user_age", "user_city")

renamed_df.show()
```

```
user_name|user_age|  user_city|
```

```
+-----+-----+-----+
```

```
|  Alice|   25|   New York|
```

```
|   Bob|   30|San Francisco|
```

```
|  Cathy|   35| Los Angeles
```

Renaming Multiple Columns

```
renamed_df = sample_df.withColumnRenamed("name", "user_name") \
    .withColumnRenamed("age", "user_age") \
    .withColumnRenamed("city", "user_city")
renamed_df.show()
```

```
user_name|user_age|  user_city|
```

```
+-----+-----+-----+
```

```
|  Alice|   25|   New York|
```

```
|   Bob|   30|San Francisco|
```

```
|  Cathy|   35| Los Angeles
```

PySpark Filter vs Where

```
import findspark
```

```
findspark.init()
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \
    .appName("Filtering Rows in PySpark DataFrames") \
    .getOrCreate()
```

```
from pyspark.sql import Row

data = [
    Row(id=1, name="Alice", age=30),
    Row(id=2, name="Bob", age=25),
    Row(id=3, name="Charlie", age=35),
    Row(id=4, name="David", age=28)]

columns = ["id", "name", "age"]

df = spark.createDataFrame(data, columns)

df.show()
```

```
id|  name|age|
+---+-----+---+
|  1|  Alice| 30|
|  2|   Bob| 25|
|  3|Charlie| 35|
|  4|  David| 28|
```

Different ways to filter rows in PySpark DataFrames

1. Filtering Rows Using 'filter' Function

It takes a boolean expression as an argument and returns a new DataFrame containing only the

rows that satisfy the condition.

Example: Filter rows with age greater than 30

```
filtered_df = df.filter(df.age > 29)
```

```
filtered_df.show()
```

```
id|  name|age|
```

```
+---+-----+---+
```

```
| 1|  Alice| 30|
```

```
| 3|Charlie| 35|
```

2. Filtering Rows Using 'where' Function

The where function is an alias for the 'filter' function and can be used interchangeably. It also takes a boolean expression as an argument and returns a new DataFrame containing only the rows that satisfy the condition.

PySpark Filter vs Where – Comprehensive Guide Filter Rows from PySpark DataFrame

Jagdeesh

Apache PySpark is a popular open-source distributed data processing engine built on top of the Apache Spark framework. It provides a high-level API for handling large-scale data processing tasks in Python, Scala, and Java.

One of the most common tasks when working with PySpark DataFrames is filtering rows based on certain conditions. In this blog post, we'll discuss different ways to filter rows in PySpark DataFrames, along with code examples for each method.

Different ways to filter rows in PySpark DataFrames

1. Filtering Rows Using 'filter' Function
2. Filtering Rows Using 'where' Function
3. Filtering Rows Using SQL Queries
4. Combining Multiple Filter Conditions

Before we dive into filtering rows, let's quickly review some basics of PySpark DataFrames. To work with PySpark DataFrames, we first need to import the necessary modules and create a `SparkSession`

```
import findspark
```

```
findspark.init()
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder \
    .appName("Filtering Rows in PySpark DataFrames") \
    .getOrCreate()
```

Next, let's create a simple DataFrame to use in our examples

```
from pyspark.sql import Row
```

```
data = [  
    Row(id=1, name="Alice", age=30),  
    Row(id=2, name="Bob", age=25),  
    Row(id=3, name="Charlie", age=35),  
    Row(id=4, name="David", age=28)
```



```
]
```

```
columns = ["id", "name", "age"]
```

```
df = spark.createDataFrame(data, columns)
```

```
df.show()
```

```
+---+-----+---+
| id|  name|age|
+---+-----+---+
|  1| Alice| 30|
|  2|  Bob| 25|
|  3|Charlie| 35|
|  4| David| 28|
+---+-----+---+
```

1. Filtering Rows Using 'filter' Function

The filter function is one of the most straightforward ways to filter rows in a PySpark DataFrame. It takes a boolean expression as an argument and returns a new DataFrame containing only the rows that satisfy the condition.

Example: Filter rows with age greater than 30

```
filtered_df = df.filter(df.age > 29)
```

```
filtered_df.show()
```

```
+---+-----+---+
```

```
| id|  name|age|
+---+-----+---+
|  1| Alice| 30|
|  3|Charlie| 35|
+---+-----+---+
```

2. Filtering Rows Using 'where' Function

The where function is an alias for the 'filter' function and can be used interchangeably. It also takes a boolean expression as an argument and returns a new DataFrame containing only the rows that satisfy the condition.

Example: Filter rows with name equal to "Alice":

```
filtered_df = df.where(df.name.isin(["Alice", "Charlie"]))
filtered_df.show()
```

```
id|  name|age|
+---+-----+---+
|  1| Alice| 30|
|  3|Charlie| 35|
```

3. Filtering Rows Using SQL Queries

Example: Filter rows with age less than or equal to 25

```
df.createOrReplaceTempView("people")
filtered_df = spark.sql("SELECT * FROM people WHERE age <= 25")
filtered_df.show()
```

```
id|name|age|
```

```
+---+---+---+
```

```
| 2| Bob| 25|
```

4. Combining Multiple Filter Conditions

combine multiple filter conditions using the '&' (and), '|' (or), and '~' (not) operators. Make sure to use parentheses to separate different conditions, as it helps maintain the correct order of operations.

```
filtered_df = df.filter((df.age > 25) & (df.name != "David"))
```

```
filtered_df.show()
```

```
id| name|age|
```

```
+---+-----+---+
```

```
| 1| Alice| 30|
```

```
| 3| Charlie| 35
```

PySpark orderBy() and sort()

```
import findspark
```

```
findspark.init()
```

```
from pyspark.sql import SparkSession
```

```
# Create a SparkSession
```

```
spark = SparkSession.builder \
```

```

.appName("PySpark orderBy() and sort() Example") \
    .getOrCreate()

# Sample data

data = [

    ("Alice", 30, "New York"),

    ("Bob", 28, "San Francisco"),

    ("Charlie", 34, "Los Angeles"),

    ("Diana", 29, "Chicago")

]

# Create a DataFrame

columns = ["Name", "Age", "City"]

df = spark.createDataFrame(data, columns)

df.show()

```

```

Name|Age|    City|
+----+---+-----+
| Alice| 30|  New York|
|  Bob| 28|San Francisco|
|Charlie| 34| Los Angeles|
| Diana| 29|   Chicago

```

orderBy() function

The **orderBy()** function in PySpark is used to sort a DataFrame based on one or more columns. It takes one or more columns as arguments and returns a new DataFrame sorted by

the specified columns.

Syntax:

```
DataFrame.orderBy(*cols, ascending=True)
```

The **sort()** Function

The **sort()** function is an alias of `orderBy()` and has the same functionality. The syntax and parameters are identical to `orderBy()`.

Syntax:

```
DataFrame.sort(*cols, ascending=True)
```

Difference between `orderBy()` and `sort()`

There is no functional difference between `orderBy()` and `sort()` in PySpark. The `sort()` function is simply an alias for `orderBy()`.

```
sorted_by_age = df.orderBy("Age")
```

```
sorted_by_age.show()
```

```
Name|Age|    City|
+-----+---+-----+
|  Bob| 28|San Francisco|
| Diana| 29|   Chicago|
| Alice| 30|  New York|
|Charlie| 34| Los Angeles|
```

```
# Sort by multiple columns using orderBy()
```

```
sorted_by_age_and_city = df.orderBy(["Age", "City"], ascending=[True, False])
```

```
sorted_by_age_and_city.show()
```

```
Name|Age|    City|
+-----+---+-----+
|  Bob| 28|San Francisco|
| Diana| 29|   Chicago|
| Alice| 30|  New York|
|Charlie| 34| Los Angeles|
```

```
# Sort the DataFrame using sort()
```

```
sorted_by_age = df.sort("Age")
```

```
sorted_by_age.show()
```

```
Name|Age|    City|
+-----+---+-----+
|  Bob| 28|San Francisco|
| Diana| 29|   Chicago|
| Alice| 30|  New York|
|Charlie| 34| Los Angeles|
```

```
# Sort by multiple columns using sort()
```

```
sorted_by_age_and_city = df.sort(["Age", "City"], ascending=[True, False])
```

```
sorted_by_age_and_city.show()
```

```
Name|Age|    City|
+-----+---+-----+
|  Bob| 28|San Francisco|
|  Diana| 29|    Chicago|
|  Alice| 30|   New York|
|Charlie| 34| Los Angeles|
```

Sorting a DataFrame using column expressions

```
from pyspark.sql.functions import desc, asc
```

```
# Sort the DataFrame by age in descending order using column expressions
```

```
sorted_by_age_desc_expr = df.orderBy(desc("Age"))
```

```
sorted_by_age_desc_expr.show()
```

```
Name|Age|    City|
+-----+---+-----+
|Charlie| 34| Los Angeles|
|  Alice| 30|   New York|
|  Diana| 29|    Chicago|
```

```
| Bob| 28|San Francisco
```

```
# Sort the DataFrame by city in ascending order using column expressions
```

```
sorted_by_city_asc_expr = df.sort(asc("City"))
```

```
sorted_by_city_asc_expr.show()
```

```
Name|Age|    City|
```

```
+-----+---+-----+
```

```
| Diana| 29|   Chicago|
```

```
|Charlie| 34| Los Angeles|
```

```
| Alice| 30|   New York|
```

```
| Bob| 28|San Francisco|
```

```
Sorting a DataFrame with NULL values
```

```
data_with_nulls = [
```

```
    ("Alice", None, "New York"),
```

```
    ("Bob", 28, None),
```

```
    ("Charlie", 34, "Los Angeles"),
```

```
    ("Diana", 29, "Chicago")
```

```
]
```

```
# Create a DataFrame with NULL values
```

```
df_with_nulls = spark.createDataFrame(data_with_nulls, columns)
```

```
# Sort the DataFrame with NULL values in Age column (NULLs appear last)
```

```
sorted_with_nulls = df_with_nulls.orderBy("Age", ascending=True, nulls_last=True)
```



```
sorted_with_nulls.show()
```

```
Name| Age|   City|
+----+----+-----+
| Alice| null| New York|
|  Bob| 28|    null|
| Diana| 29|  Chicago|
|Charlie| 34|Los Angeles
```

```
# Sort the DataFrame with NULL values in City column (NULLs appear first)
sorted_with_nulls_alt = df_with_nulls.sort("City", ascending=True, nulls_first=True)
sorted_with_nulls_alt.show()
```

```
Name| Age|   City|
+----+----+-----+
|  Bob| 28|    null|
| Diana| 29|  Chicago|
|Charlie| 34|Los Angeles|
| Alice| null| New York
```

Sorting a DataFrame using a custom sorting order

```
from pyspark.sql.functions import col, when

# Define a custom sorting order for cities
city_order = ["New York", "Los Angeles", "Chicago", "San Francisco"]
```

```
# Create a custom sorting column

custom_sort_col = when(col("City") == city_order[0], 0) \

    .when(col("City") == city_order[1], 1) \

    .when(col("City") == city_order[2], 2) \

    .when(col("City") == city_order[3], 3) \

    .otherwise(4)

# Sort the DataFrame using the custom sorting order

sorted_by_custom_order = df.orderBy(custom_sort_col)

sorted_by_custom_order.show()

Name|Age|    City|
+----+---+-----+
| Alice| 30| New York|
|Charlie| 34| Los Angeles|
| Diana| 29| Chicago|
| Bob| 28| San Francisco
```

PySpark GroupBy()

PySpark GroupBy is a powerful operation that allows you to perform aggregations on your data. It groups the rows of a DataFrame based on one or more columns and then applies an aggregation function to each group. Common aggregation functions include sum, count, mean, min, and max.

Here's a general structure of a GroupBy operation:

Syntax :

```
dataFrame.groupBy("column_name").agg(aggregation_function)
```

aggregation functions

count() – return the number of rows for each group

max() – returns the maximum of values for each group

min() – returns the minimum of values for each group

sum() – returns the total for values for each group

avg() – returns the average for values for each group

```
import findspark
```

```
findspark.init()
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import *
```

```
# Create a Spark session
```

```
spark = SparkSession.builder \
    .appName("PySpark GroupBy Example") \
    .getOrCreate()
```

```
# Sample data
```

```
data = [("1001", "Laptop", "Electronics", 1, 1000, "2023-01-01"),
        ("1002", "Mouse", "Electronics", 2, 50, "2023-01-02"),
        ("1003", "Laptop", "Electronics", 1, 1200, "2023-01-03"),
        ("1004", "Mouse", "Electronics", 3, 30, "2023-01-04"),
        ("1005", "Smartphone", "Electronics", 1, 700, "2023-01-05")]
```

```
# Create DataFrame
```

```
columns = ["OrderID", "Product", "Category", "Quantity", "Price", "Date"]
```

```
df = spark.createDataFrame(data, columns)
```

```
df.show()
```

```
OrderID| Product| Category|Quantity|Price| Date|
+-----+-----+-----+-----+-----+-----+
| 1001| Laptop|Electronics| 1| 1000|2023-01-01|
| 1002| Mouse|Electronics| 2| 50|2023-01-02|
| 1003| Laptop|Electronics| 1| 1200|2023-01-03|
| 1004| Mouse|Electronics| 3| 30|2023-01-04|
| 1005|Smartphone|Electronics| 1| 700|2023-01-05
```

GroupBy operation on single column

```
# GroupBy and aggregate
```

```
result = df.groupBy("Product").agg(sum("Price").alias("Total_Sales"))
```

```
# Show results
```

```
result.show()
```

```
Product|Total_Sales|
+-----+-----+
| Laptop|    2200|
| Mouse|     80|
|Smartphone|    700
```

GroupBy operation on Multiple Columns

```
# GroupBy and aggregate

result = df.groupby(["Product", "Category"]) \

    .agg(sum("Price").alias("Total_Sales"))

# Show results

result.show()
```

```
Product|  Category|Total_Sales|
+-----+-----+-----+
|  Laptop|Electronics|    2200|
|  Mouse|Electronics|     80|
|Smartphone|Electronics|    700|
```

GroupBy operation on Multiple Aggregations

```
# GroupBy and aggregate

result = df.groupby("Product") \

    .agg(sum("Price").alias("Total_Sales"),

        sum("Quantity").alias("Total_Quantity"))

# Show results

result.show()
```

```
Product|Total_Sales|Total_Quantity|
+-----+-----+-----+
|  Laptop|    2200|         2|
```

Mouse	80	5
Smartphone	700	1

er Aggregated data using where condition

can use a combination of where() (which is equivalent to the SQL WHERE clause) and groupBy() to perform a groupBy operation with a specific condition.

GroupBy and aggregate using where condition

```
result = df.groupBy("Product") \
    .agg(avg("Price").alias("Total_Sales"),
         sum("Quantity").alias("Total_Quantity")) \
    .where(col("Total_Quantity") >= 2)
```

Show results

```
result.show()
```

Product	Total_Sales	Total_Quantity
---------	-------------	----------------

+-----+	+-----+	+-----+
---------	---------	---------

Laptop	1100.0	2
--------	--------	---

Mouse	40.0	5
-------	------	---

Custom Aggregation Functions

```
import pandas as pd
```

```
from pyspark.sql.types import FloatType
```

```

from pyspark.sql.functions import pandas_udf

@pandas_udf(FloatType())

def median(column: pd.Series) -> float:

    return float(column.median())

```

Category | Median_Price |

+-----+-----+

|Electronics| 500.0

PySpark Joins

Type of Joins

Inner Join

Outer (Full) Join

Left Join

Right Join

Left Semi Join

Left Anti Join

Cross Join

```
import findspark
```

```
findspark.init()
```

```
from pyspark.sql import SparkSession
```

```
# Initialize Spark Session
```

```
spark = SparkSession.builder.master("local").appName("PySpark Join Types").getOrCreate()
```

```
# Create sample dataframes
```

```
df1 = spark.createDataFrame([(1, "A"), (2, "B"), (3, "C")], ["id", "value1"])
```

```
df2 = spark.createDataFrame([(1, "X"), (2, "Y"), (4, "Z")], ["id", "value2"])
```

```
# Perform inner join
```

```
result = df1.join(df2, on="id", how="inner")
```

```
# Show result
```

```
result.show()
```

```
id|value1|value2|
```

```
+---+-----+-----+
```

```
| 1|  A|  X|
```

```
| 2|  B|  Y|
```

```
# Perform outer join
```

```
result = df1.join(df2, on="id", how="outer")
```

```
# Show result
```

```
result.show()
```

```
id|value1|value2|
```

```
+---+-----+-----+
```



```
| 1|  A|  X|
| 2|  B|  Y|
| 3|  C| null|
| 4| null|  Z
```

```
# Perform left join
```

```
result = df1.join(df2, on="id", how="left")
```

```
# Show result
```

```
result.show()
```

```
id|value1|value2|
```

```
+---+-----+-----+
```

```
| 1|  A|  X|
```

```
| 3|  C| null|
```

```
| 2|  B|  Y|
```

```
# Perform right join
```

```
result = df1.join(df2, on="id", how="right")
```

```
# Show result
```

```
result.show()
```

```
+---+-----+-----+
```

```
| id|value1|value2|
```

```
+---+-----+
| 1|  A|  X|
| 2|  B|  Y|
| 4| null|  Z|
```

Left Semi Join

A left semi join returns only the columns from the left dataframe for the rows with matching keys in both dataframes. It is similar to an inner join but only returns the columns from the left dataframe.

```
# Perform left semi join
```

```
result = df1.join(df2, on="id", how="left_semi")
```

```
# Show result
```

```
result.show()
```

```
+---+-----+
| id|value1|
+---+-----+
| 1|  A|
| 2|  B|
```

Left Anti Join

A left anti join returns the rows from the left dataframe that do not have matching keys in the right dataframe. It is the opposite of a left semi join.

```
# Perform left anti join
```

```
result = df1.join(df2, on="id", how="left_anti")
```

```
# Show result
```

```
result.show()
```

```
+---+-----+
```

```
| id|value1|
```

```
+---+-----+
```

```
| 3|    C|
```

Cross Join

A cross join, also known as a cartesian join, returns the cartesian product of both dataframes. It combines each row from the left dataframe with each row from the right dataframe.

```
# Perform cross join
```

```
result = df1.crossJoin(df2)
```

```
# Show result
```

```
result.show()
```

```
+---+-----+---+-----+
```

```
| id|value1| id|value2|
```

```
+---+-----+---+-----+
```

```
| 1|    A| 1|    X|
```

```
| 1|    A| 2|    Y|
```

```
| 1|    A| 4|    Z|
```

```
| 2|  B| 1|  X|
| 2|  B| 2|  Y|
| 2|  B| 4|  Z|
| 3|  C| 1|  X|
| 3|  C| 2|  Y|
| 3|  C| 4|  Z|
```

PySpark Union?

PySpark Union is an operation that allows you to combine two or more DataFrames with the same schema, creating a single DataFrame containing all rows from the input DataFrames.

It's important to note that the Union operation doesn't eliminate duplicate rows, so you may need to use the `distinct()` function afterward if you want to remove duplicates.

```
import findspark

findspark.init()

from pyspark.sql import SparkSession

from pyspark.sql.types import StructType, StructField, StringType, IntegerType

# Create a Spark session

spark = SparkSession.builder.appName("PySpark Union Example").getOrCreate()

# Define the schema

schema = StructType([
```

```

    StructField("product", StringType(), True),
    StructField("price", IntegerType(), True),
    StructField("quantity", IntegerType(), True)
])

# Create DataFrame for region A
data_A = [("apple", 3, 5), ("banana", 1, 10), ("orange", 2, 8)]
df_A = spark.createDataFrame(data_A, schema=schema)

# Create DataFrame for region B
data_B = [("apple", 3, 5), ("banana", 1, 15), ("grape", 4, 6)]
df_B = spark.createDataFrame(data_B, schema=schema)

# Create DataFrame for region C
data_C = [("apple", 3, 10), ("banana", 1, 20), ("grape", 4, 10), ("orange", 2, 7)]
df_C = spark.createDataFrame(data_C, schema=schema)

# Perform the Union operation on two DataFrames
df_union = df_A.union(df_B)

# Show the results
df_union.show()

+-----+-----+-----+
|product|price|quantity|
+-----+-----+-----+
| apple|   3|     5|
| banana|  1|    10|

```

```
| orange|  2|   8|
| apple|  3|   5|
| banana| 1|  15|
| grape|  4|   6|
```

Union without Duplicates

It's important to note that the Union operation doesn't eliminate duplicate rows, so you may need to use the `distinct()` function afterward if you want to remove duplicates.

```
# Perform the Union operation on two DataFrames
```

```
df_union_dist = df_A.union(df_B).distinct()
```

```
# Show the results
```

```
df_union_dist.show()
```

```
product|price|quantity|
```

```
+-----+-----+-----+
```

```
| apple|  3|   5|
```

```
| banana| 1|  10|
```

```
| orange|  2|   8|
```

```
| banana|  1|  15|
```

```
| grape|  4|   6|
```

Union Multiple DataFrames

```
# Perform the Union operation on multiple DataFrames
```

```
df_union_all = df_A.union(df_B).union(df_C)
```

```
# Show the results
```

```
df_union_all.show()
```

```
product|price|quantity|
```

```
+-----+-----+-----+
```

```
| apple| 3| 5|
```

```
| banana| 1| 10|
```

```
| orange| 2| 8|
```

```
| apple| 3| 5|
```

```
| banana| 1| 15|
```

```
| grape| 4| 6|
```

```
| apple| 3| 10|
```

```
| banana| 1| 20|
```

```
| grape| 4| 10|
```

```
| orange| 2| 7|
```